

Description Field — Generation Rules (Developer Spec, Final)

This governs how the **Description** field (Contentful metadata tab) is generated. It doubles as the meta description shown in search results — no disclaimer sits next to it, so it has to be accurate and compliant standing completely on its own.

This version supersedes any earlier draft. Notable change from the original skill: there is no “reuse existing approved copy” tier anymore — every description is freshly generated, every time, using judgment over the full article.

1. Process

Run in this order:

1. **Gate on series type.** This determines *what* gets read in step 2 — see section 2.
2. **Read the source content** (full article body, or just the top paragraph — depends on series type).
3. **Analyze and draft the summary/description**, always freshly generated — there is no external-copy shortcut to check first.
4. **Improve for SEO/AI** against the concrete rubric in section 5.
5. **Run the compliance-safety screen — last, and it overrides everything before it.** This step can force a rewrite of anything steps 3-4 produced, regardless of what it costs the SEO score. Compliance is the top priority in this process, full stop.

2. What gets read (series-specific gate)

- **Point of View articles:** check for the on-page summary blurb first — see section 2a below. Takes priority over the general full-length-piece rule.
- **Full-length pieces** (Monthly Commentary, Research papers, Trip Notes, Economic Update, etc. — Point of View only when no on-page summary is found): read and analyze the **entire article body** before drafting. Don’t generate from a headline, slug, or a single isolated paragraph in isolation.
- **Week in Review:** read **only the top/opening narrative paragraph**. “Highlights of the Week” is standalone information — it recurs with the same fixed categories every issue and is never pulled into or blended with the description, no exceptions.
- **Hub/landing pages:** no single body to read — build the description from the linked articles’ own leading title portions instead (see section 6).

Title/subtitle inform what the piece is about but are not source material for the actual wording — the description must not restate the Title (see section 5’s “Distinct from Title” rubric).

2a. Point of View — on-page summary carve-out (confirmed rule, 2026-08-19)

Applies only to Point of View articles. Many Point of View pieces already carry a short editorial summary blurb directly on the page — it sits beneath the headline, above the “Back to Library” / Table of Contents block, and before the article’s actual body text begins.

This blurb is distinct from two other things on the same page that can look similar but aren’t it: the body’s own opening paragraph, and the meta description tag. Confirmed case (2026-08-10, “Wisdom Of The Crowd”): all three texts on that page were different from each other. Don’t trust a single automated fetch to locate this blurb correctly — a first-pass fetch on this exact piece returned a fabricated summary that didn’t match the real page at all. Verify against the actual rendered page before treating any candidate text as this summary.

If the summary exists: use it directly as the Description. Don’t rewrite or paraphrase it — it’s already-approved editorial copy, and the entire point of this rule is reusing it, not generating new copy. - Under 130 characters: use as-is. - Over 130 characters (the common case — this blurb is written for the page, not for a meta-description budget): truncate at the last full word boundary. Do **not** run it through the compress-and-rewrite step described in section 4 — compressing/rewriting an official editorial summary to shorten it defeats the purpose of pulling it verbatim. - Still run the compliance-safety screen (section 3) over the result as a final check, even though this text is already-published editorial copy.

If no summary exists for a given Point of View piece, fall back to the standard full-length-piece process (read the entire body → draft → SEO → compliance).

This carve-out is specific to Point of View only — it does not change the process for Week in Review, Monthly Commentary, Trip Notes, Economic Update, or any other series.

3. If the full source can’t be reached

Fail closed. If a streaming-payload or lazy-loaded-iframe blind spot blocks the article body (or a chart/data visualization is embedded via iframe and doesn’t load), **do not guess or fabricate to fill the gap** — flag it and let the user decide how to proceed. This includes chart-only embeds: since the description is meant to stay at a general-perspective level rather than cite specific figures anyway (see section 4), there’s rarely a need to chase down an iframed chart’s contents — but if something fails to load, flag it rather than silently deciding it doesn’t matter.

4. Performance figures — restricted, not banned outright

A specific performance/return/yield figure may appear in the description **only if it is already written verbatim in the article’s own text** — never inferred, calculated, estimated, or sourced from outside the piece. Even when available, this is **not the default move**: the description should generally stay at a general/qualitative level. Use a specific number only when it’s clearly the single substantiated data point that matters — not as a routine way to sound more specific.

5. Compliance-safety screen — hard override, final gate

This is the last step in the process (section 1, step 5) and has veto power over the SEO pass. Screen against:

- No performance guarantees or promises (“will outperform,” “guaranteed returns,” “will deliver X%”).
- No absolute/certainty language about markets or outcomes (“always,” “never,” “certainly,” bare “will”).
- No individualized investment directives aimed at the reader (“investors need to,” “investors should,” “you should buy/sell/rotate into”).
- No unqualified superlatives or rankings (“the best fund,” “top-performing strategy”) unless directly attributed to a named, cited source.
- Forward-looking statements get hedged (“may,” “could,” “is likely to”) rather than stated as settled fact.
- Past performance isn’t framed as predictive of future results.
- Favor attributive framing (“discusses,” “shares her view on,” “answers questions on”) over declarative market calls (“bonds are back”).
- No bare performance figures beyond what section 4 allows.

If any phrase trips one of these, rewrite toward the safer/hedged version automatically — don’t weigh it against how well the bolder version scores on SEO. This is a drafting-level screen, not a compliance sign-off; anything genuinely borderline still gets flagged to a human rather than resolved silently.

6. Tone/voice

Attributive, professional, concise — matching the same voice the compliance framing already implies (“discusses,” “shares her view on,” “answers questions on”). No separate style system beyond this.

7. Character cap — 130 max, hard rule, compress before truncating

Check .length programmatically, not by eye. **Default to compressing/rewriting the draft to fit under 130 characters first.** Only fall back to truncating at the last full word boundary when a rewrite genuinely can’t get there — straight truncation risks cutting off the actual point rather than a filler clause, especially on a full-piece summary. Never truncate mid-word, never use an ellipsis that eats into the budget. (Google itself typically truncates around 155-160 characters — 130 is a deliberately conservative internal standard, not a platform limit.)

8. Recurring-series rule: lead with substance, not a repeated preamble

For anything on a fixed publishing cadence (weekly, monthly, etc.), never open with the same fixed phrase every time (e.g. “Payden’s Week in Review looks at...” verbatim, week after week) — search engines are likely to discard a boilerplate pattern at scale and generate their own snippet instead, meaning the compliance-reviewed copy never shows up in results. Open with the piece’s specific content instead; the series identity already lives in the Title field.

9. SEO/AI self-check (internal heuristic, not a certified score)

Score out of 100, four equally-weighted categories — an internal consistency check, never presented to the user as an actual SEO tool's output:

- **Length utilization (25 pts):** 110-130 characters = 25; 90-109 = 18; 70-89 = 10; under 70 (or only short because compression/truncation cut it down) = 5.
- **Front-loaded specificity (25 pts):** a concrete named entity/event/figure within the first ~60 characters = 25; present but after character 60 = 12; purely generic opening = 0.
- **Keyword alignment (25 pts):** a top-3 ranked keyword appears verbatim/near-verbatim = 25; a related but non-matching term = 12; no overlap = 0.
- **Distinct from Title (25 pts):** description adds information beyond the Title with no verbatim phrase lifted from it (brand/series name exempt) = 25; partial restatement = 12; mostly a rephrase of the Title = 0.

Bands: 90-100 ship as-is. 70-89 acceptable, note the weakest category. Below 70, revise once and re-score — don't loop chasing 100. Section 5's compliance screen always overrides this score regardless of how high it lands.

10. Hub/landing page sourcing

No single lead paragraph exists (it's a table of contents, not a narrative) — build the description from the linked articles' own leading title portions: Payden's {Series Name} Vol. {N}, {Year}: {Title 1}, {Title 2}, {Title 3}, {Title 4}. Same 130-character cap and full compliance screen still apply.

Summary of what changed from the original skill

- Removed the "existing approved external copy" tier entirely — there is no outside reference to check; every description is freshly generated.
- Added an explicit full-context-read requirement for full-length pieces (previously implied by Step 1's fetch, never stated as a Description-specific requirement).
- Made Week in Review's "top paragraph only, never Highlights" rule absolute — removed the prior fallback that allowed pulling from Highlights when the opening was thin.
- Added a fail-closed rule for unreachable content (including iframed charts), matching the schema step's existing disclosure-gate pattern.
- Added an explicit rule on bare performance figures (verbatim-in-source only, general framing preferred by default).
- Reordered compliance to run last, as a final override step, rather than as a mid-process check.
- Added a compress-before-truncate default for the 130-character cap.
- Added an explicit tone/voice default (attributive, professional, concise).
- **New (2026-08-19):** Point of View articles now check for an existing on-page summary blurb first (section 2a) and use it verbatim, truncating rather than compressing if it's over 130 characters — a carve-out from the general "always freshly generated" and "compress before truncate" rules, scoped only to Point of View.